Ramped AI

30-DAY GO-LIVE GUARANTEE

BUSINESS STRATEGY · COMPETITIVE ANALYSIS

Your AI department, live in 30 days.

Strategy review v2 · 2026-04-29

Audited at commit c87ffa1 · post-deploy of Phase 1+2+3

THE STRATEGIC FRAME

Ramped sells the only thing AI tools can't: the certainty that someone else owns the outcome.

Build everything else around that.

WHAT RAMPED SELLS TODAY

Three SKUs. One promise. Mid-market done-for-you.

Starter

$2,500/mo

$2,083/mo billed annually · save $5,000

+ $2,500 onboarding

- 1–2 agents
- Build → live in 30 days
- Slack support
- Weekly digest

MOST POPULAR

Growth

$5,000/mo

$4,167/mo billed annually · save $10,000

+ $3,500 onboarding

- 3–5 agents
- Build → live in 30 days
- Quarterly business review
- Approvals queue + portal access

Enterprise

From $10K/mo

Scoped on call · custom SLA

Onboarding scoped

- Unlimited agents
- Dedicated implementation lead
- On-prem / VPC option (roadmap)
- Compliance: SOC 2 in progress

30-day go-live guarantee — full refund if your first agent isn't live by day 30. No fine print. No partial payments.

UNIT ECONOMICS

Healthy cash profile. The leverage is in templating.

$53,500

Annual ACV (Growth annual)

12 mo prepaid + onboarding · paid in 14 days

50–70%

Gross margin (templated)

drops to 25–40% if every customer is bespoke

30 days

Time-to-deliver (the promise)

hard ceiling · drives ops cadence

≈14 days

Time-to-cash (annual)

from booking to invoice paid

21% / yr

AI-app retention benchmark

vs 31% non-AI · done-for-you must beat materially

<5%

Refund rate target

tail-risk metric · publish once stable

Templated reuse > headcount: the 4th deployment of "Inbound Quote Drafter" should take ¼ the time of the 1st.

COMPETITIVE LANDSCAPE

Done-for-you vs self-serve · 10 players + 2 substrate threats

Competitor	Model	Floor price	Risk reversal	Vector
Ramped AI	Done-for-you · retainer	$2,500/mo	30-day refund ✓ unique	—
Lindy.ai	Self-serve · platform	Free / $50	None	DIY substitute · 1/12 the price
Gumloop	Self-serve · workflow builder	Free / $37	None	DIY · prosumer drag-drop
Zapier AI Agents	Self-serve · add-on	$33+ on top	None	DIY for existing Zapier users
Viktor	Self-serve · Slack-native	Free / $50	None	Closest positioning rival; cheaper
WhiteHorse AI	Done-for-you · perf-priced	Per appt.	Pay only when it works	Strongest risk reversal · sales only
GoHighLevel AI	Agency-resold · platform	$97 + $97/mo	None	Indirect · vertical via agencies
Adept AI	(Amazon-acquired Jun 2024)	—	—	Defunct as standalone
SmythOS	Self-serve · on-prem option	$30/mo	None	Adjacent · dev-focused
Automation Agency	Done-for-you · low-end	$249/mo	30-day guarantee	Indirect · mass-market
AI SMB Solutions	Done-for-you · consulting	Custom	10–20 hrs claim	Same promise · less productized
Claude Cowork / ChatGPT Agent	Substrate · self-serve	$20–200/mo	—	Substrate threat — they sell software, you sell outcome

Pricing band: Ramped sits above every self-serve tool ($30–200/mo) and below custom AI engagements ($30K–100K upfront).

POSITION VS THE FIELD

The done-for-you / mid-market quadrant is empty except Ramped.

⚠ *Claude Cowork + ChatGPT Agent commoditize the substrate. Compete on outcome + accountability, not the agent itself.*

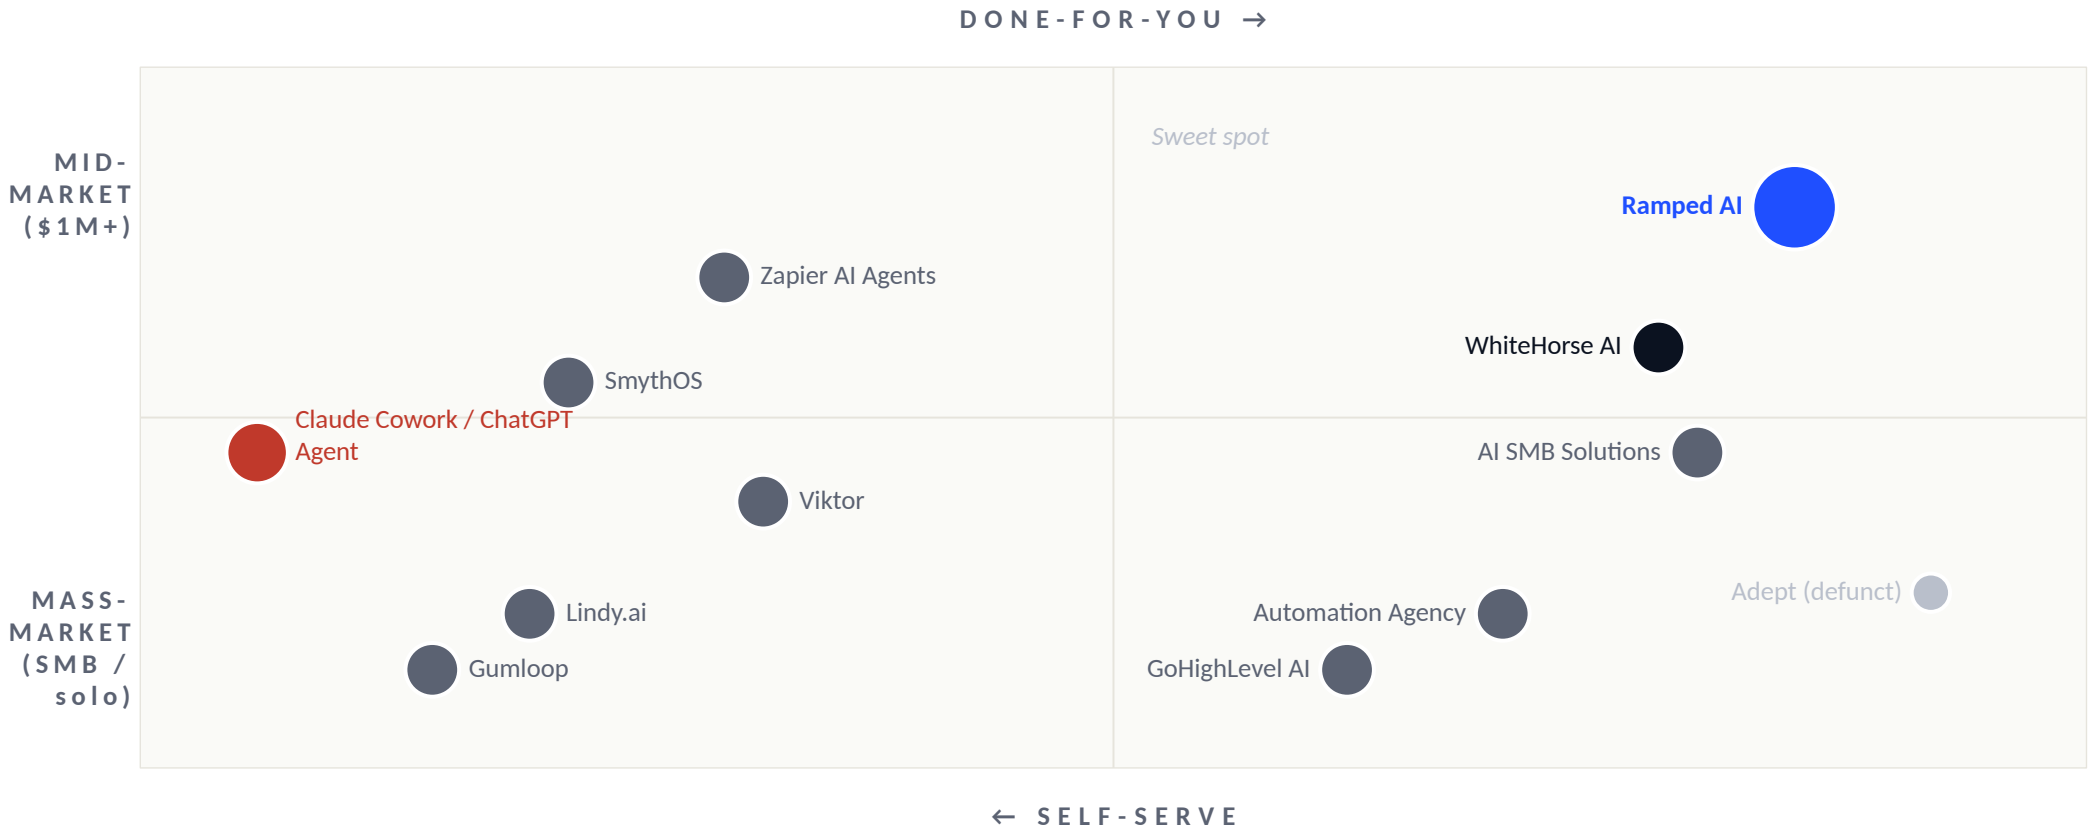

What's defensible — keep doing this.

1 Time-bound + money-back guarantee

30-day refund is the only falsifiable risk reversal in the matrix. Closest analogue: Stripe's 'no lost checkout' guarantee.

2 Visible, transparent pricing

Among done-for-you peers, Ramped is the only one with public monthly tier pricing. Removes a discovery-call friction step.

3 Sharp product surface

Customer portal (post-Phase 1+2 hardening), phase indicator, draft approvals, weekly digest. Retention engineering most $5K/mo agencies don't have.

4 Operator-narrative founder

Andrew Yoon (10-yr operator at Xtractor Depot, supplied SpaceX/Lucid Motors) matches the buyer (CEO/operator) better than 'ex-FAANG engineer.'

5 Strong audit + delivery muscle

v1+v2 audits, Playwright suite, RLS-hardened DB, signed-token auth. Compounds. Most $5K/mo agencies don't have this.

What's not defensible yet — fix in priority order.

1 Single-customer social proof

One quote (Xtractor Depot) with three 'Verified' chips. At $30K+ ACV, sharp CFOs ask for references.

2 Vertical claims aren't backed

8 industries pilled; 1 shipped (industrial supply). Risk if a prospect's industry is on the list.

3 No live demo

Phase 2 V2-B6 added a CSS overlay. Real video or sandbox would 2–3× hero conversion.

4 'AI department' is unowned naming

Lindy: 'AI assistants.' Viktor: 'AI coworkers.' HighLevel: 'AI Employees.' Wrap it in something proprietary by Q3.

5 Refund-rate is binary

One public refund tweet erodes the central commitment. Operationally the #1 metric to watch.

6 Customer portal is read-mostly

Customers approve drafts; can't configure agents. Indispensability ↑ retention, but bad SaaS-multiple optics.

Where the next 12 months of growth lives.

1 Vertical specialization

Pick 2 verticals (industrial supply + dental/field svc). Per-vertical landing page should triple conversion vs generic /book.

2 Productized add-ons

$497 'AI Health Check' audit, free 'starter agent recipe' library. Bridges the $0–$2,500 price gap.

3 Second-customer flywheel

Use next 3 wins for named testimonials, photographed logos, ≥1 video. Replace placeholder customer-strip logos as they ship.

4 'Ramp Index' report

Quarterly anonymized benchmark report. Inbound SEO + earned media engine. Wait until n≥10.

5 White-label channel

License playbook to a regional MSP / accounting firm. Low CAC, high margin. Wait for n=10 + senior IM.

6 '0 of N refunded' counter

Once true and ratio stable, publish prominently. A marketing weapon Lindy can never match.

Watch closely. Two are existential.

1 Substrate commoditization

Claude Cowork + ChatGPT Agent + OpenAI Codex Background Computer Use launched 2025–26 at $20–200/mo. They sell software; you sell outcomes + accountability. Re-emphasize human-managed angle.

2 DIY tools narrow the gap

Lindy + Gumloop visibly better month-over-month. Buyer who needed Ramped 12 mo ago can now spend 4 hrs in Lindy. Question: do they want to spend those 4 hrs?

3 AI-app retention crisis

Industry: AI products retain 30% worse than non-AI year-1 (21% vs 31% annual). Done-for-you is best structural defense — but only if outcome metrics are visible monthly.

4 Founder bandwidth cliff

Jon does every discovery call. ~6–8 calls/week × 25% close = 8 customers/month max. Hire #1 = senior implementation manager by end of Q2.

PRICING TESTS TO RUN

What works. What I'd test. What I would NOT change.

Test	Hypothesis	Effort
Add a "Try It" tier at $497/mo	One agent, no onboarding, deployed in 14 days. Catches buyers who can't justify $2.5K. Upsell to Starter at month 3.	High
Move Starter to $1,997/mo monthly	Reads materially below the psychological $2K barrier. Marginal revenue loss; meaningful conversion lift.	Low
Bundle onboarding into first month for annual	Annual buyers commit — don't make them feel double-charged. "$25K/yr includes setup" is cleaner.	Low
Add a Pro tier at $3,500/mo	$2.5K → $5K is steep. Middle tier captures Starter-outgrowers without forcing Growth's volume.	Medium
Post Enterprise floor: "from $10K/mo"	"Custom" reads as "we'll figure it out." A floor primes anchoring + signals seriousness.	Low
Refund-rate counter (when ≥30 customers · 0)	Public commitment-device evidence. Marketing weapon Lindy can't match.	Low (data-gated)

DO NOT CHANGE: the $5K/mo Growth ceiling · the 30-day guarantee window (don't extend or shorten) · the flat-monthly framing (per-seat fights the 'you don't manage it' message).

POSITIONING GAPS

Site-level fixes mapped to competitor moves.

Gap	Today	Fix
"Why not Lindy / Zapier / Gumloop?"	Comparison covers VA-vs-AI only	Add Tab 2 to /comparison: 3 reasons buyers pay 30× more for done-for-you
Compliance / security signals	None visible	Add SOC 2 (in progress) + sub-processors to footer; start Type 1 by Q3
Vertical proof	8 industries listed; 1 shipped	Pick 2 verticals; build 1-page landing each; drop unbacked industries
Refund commitment as a number	Stated as policy	Once data supports it: "Refunded 0 of N customers since [date]" hero stat
Self-serve consideration cost	$2,500 floor; no free trial	Add "$0 — see what we'd build" personalized roadmap as marketing asset
Founder credibility above-fold	Andrew Yoon line present ✓	Verify SpaceX/Lucid claim; add proof line if real, soften if not
Quarterly Business Review	Not advertised	Make QBR a Growth-tier benefit. Frames renewal around outcomes, not invoice

ACTION PLAN

Sequenced. Reversible. Compound.

THIS WEEK

5 items · <3 hrs each

1. Add Tab 2 to /comparison: vs DIY tools (Lindy, Zapier, Gumloop)
2. Trim industry pill list 8 → 2–3 verticals you can back
3. Promote 30-day guarantee from chip → hero co-equal
4. Add Enterprise floor: "from $10K/mo · scoped on call"
5. Verify SpaceX/Lucid claim once and for all

NEXT 30 DAYS

5 strategic moves

1. Land 2 named testimonials (target 2026-05-29)
2. Build 1 vertical landing page (e.g. /industrial-supply)
3. Add "real demo" video on hero (60-sec Loom or screencast)
4. Productize questionnaire as standalone lead magnet
5. Start SOC 2 Type 1 (Vanta / Drata)

NEXT QUARTER

5 strategic moves

1. ★ Hire #1: Senior Implementation Manager
2. Build 2 vertical playbooks (templated build process)
3. Launch "Ramp Index" benchmark report (post n≥10)
4. Refund-rate dashboard (internal → public when stable)
5. Add $497 "AI Health Check" SKU

★ Single highest-leverage move: hire #1 unblocks the founder bottleneck. Without it, Ramped is capped at ~10 active customers.

METRICS DASHBOARD

10-row weekly view. Most fields already in Supabase.

Metric	Why it matters	Source
Live customer count	Monthly progression · vanity but motivating	bookings.payment_status = subscription_active
30-day on-time delivery rate	The brand-defining commitment	bookings.go_live_at ≤ datetime + 30d
Refund rate	Tail-risk metric · brand safety	Stripe events / new bookings.refunded_at column
Avg time-to-first-agent-live	Build velocity; should improve as playbooks land	agent.created_at + first agent_runs
Hours saved / customer / week	Customer-success metric · the renewal narrative	aggregate from agent_runs.hours_saved
Discovery-call → booking conv.	Top-funnel health	Stripe-paid / booked discovery calls
Booking → questionnaire-completed	Mid-funnel · if low, questionnaire is friction	bookings with non-null questionnaire
Roadmap-emailed → booking-paid	Late funnel · conversion-killer if low	automation_map field + Stripe events
Net Revenue Retention (NRR)	Upgrade $ vs churn $ · SaaS-grade health	Stripe subscriptions + invoice history
Refund $ / total revenue	One Growth-annual refund = 10× Starter monthly	Stripe

RISK REGISTER

Top 5 in priority order.

1	**Refund-rate cliff**	One refund early matters more than 10 in steady state. Aggressive scoping on discovery > forcing a marginal customer through onboarding.
2	**Substrate commoditization**	Anthropic Cowork + OpenAI Agent target the substrate. Compete on outcome + accountability, not the agent itself.
3	**Founder bandwidth cap**	Jon's 6–8 calls/week × close-rate caps revenue at ~10 active customers. Unblock with hire #1 by end of Q2.
4	**First failed delivery**	SMB referral networks tight; one missed 30-day promise lands in a Slack thread. Ship a 'war room' framework for any deal at risk by day 21.
5	**Pricing-toggle copy drift**	v1 audit found 3 places showed different prices. Re-occurs on every price change. Extend lint scripts to check pricing consistency.

THE FRAME

Ramped sells the only thing AI tools can't: the certainty that someone else owns the outcome.

- 30-day guarantee
- Done-for-you
- Human-managed

Strategy review v2 · 2026-04-29 · See BUSINESS-STRATEGY.md for the full doc.